

## Chapter 4

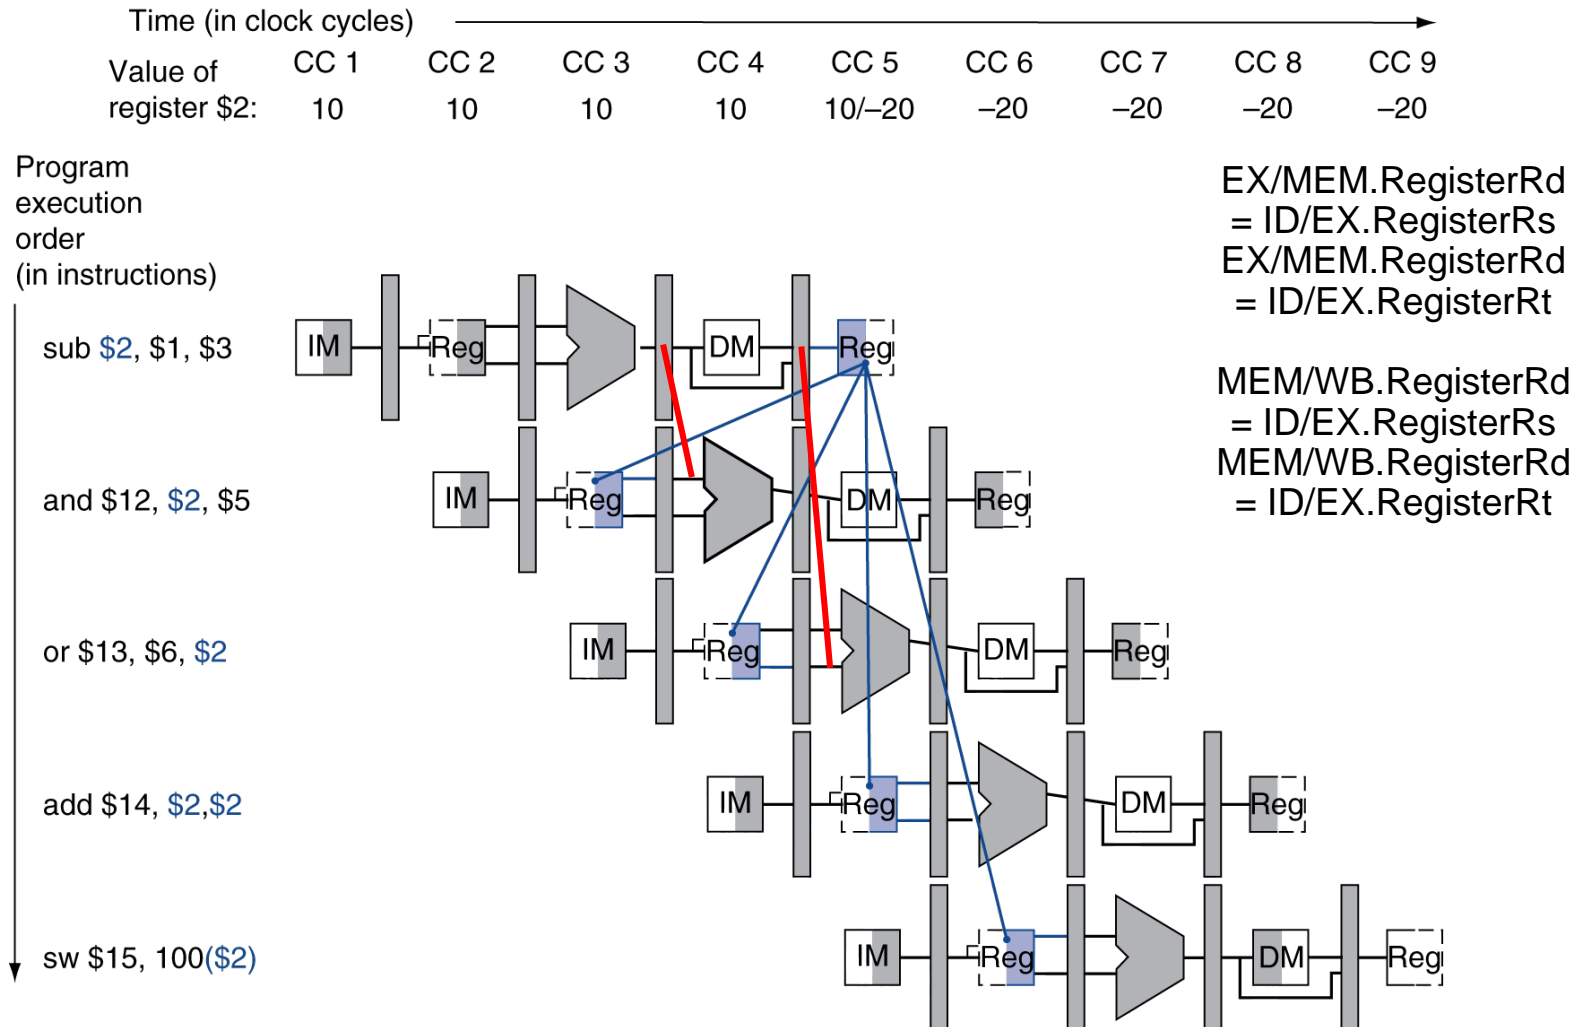
Part D :

Datapath and Control for Data and  
Control Hazards

Exceptions

Parallelism via Instructions

# Data Hazard Example



# Detecting the Need to Forward

- Data hazards when

1a. EX/MEM.RegisterRd = ID/EX.RegisterRs

1b. EX/MEM.RegisterRd = ID/EX.RegisterRt

2a. MEM/WB.RegisterRd = ID/EX.RegisterRs

2b. MEM/WB.RegisterRd = ID/EX.RegisterRt

Fwd from  
EX/MEM  
pipeline reg

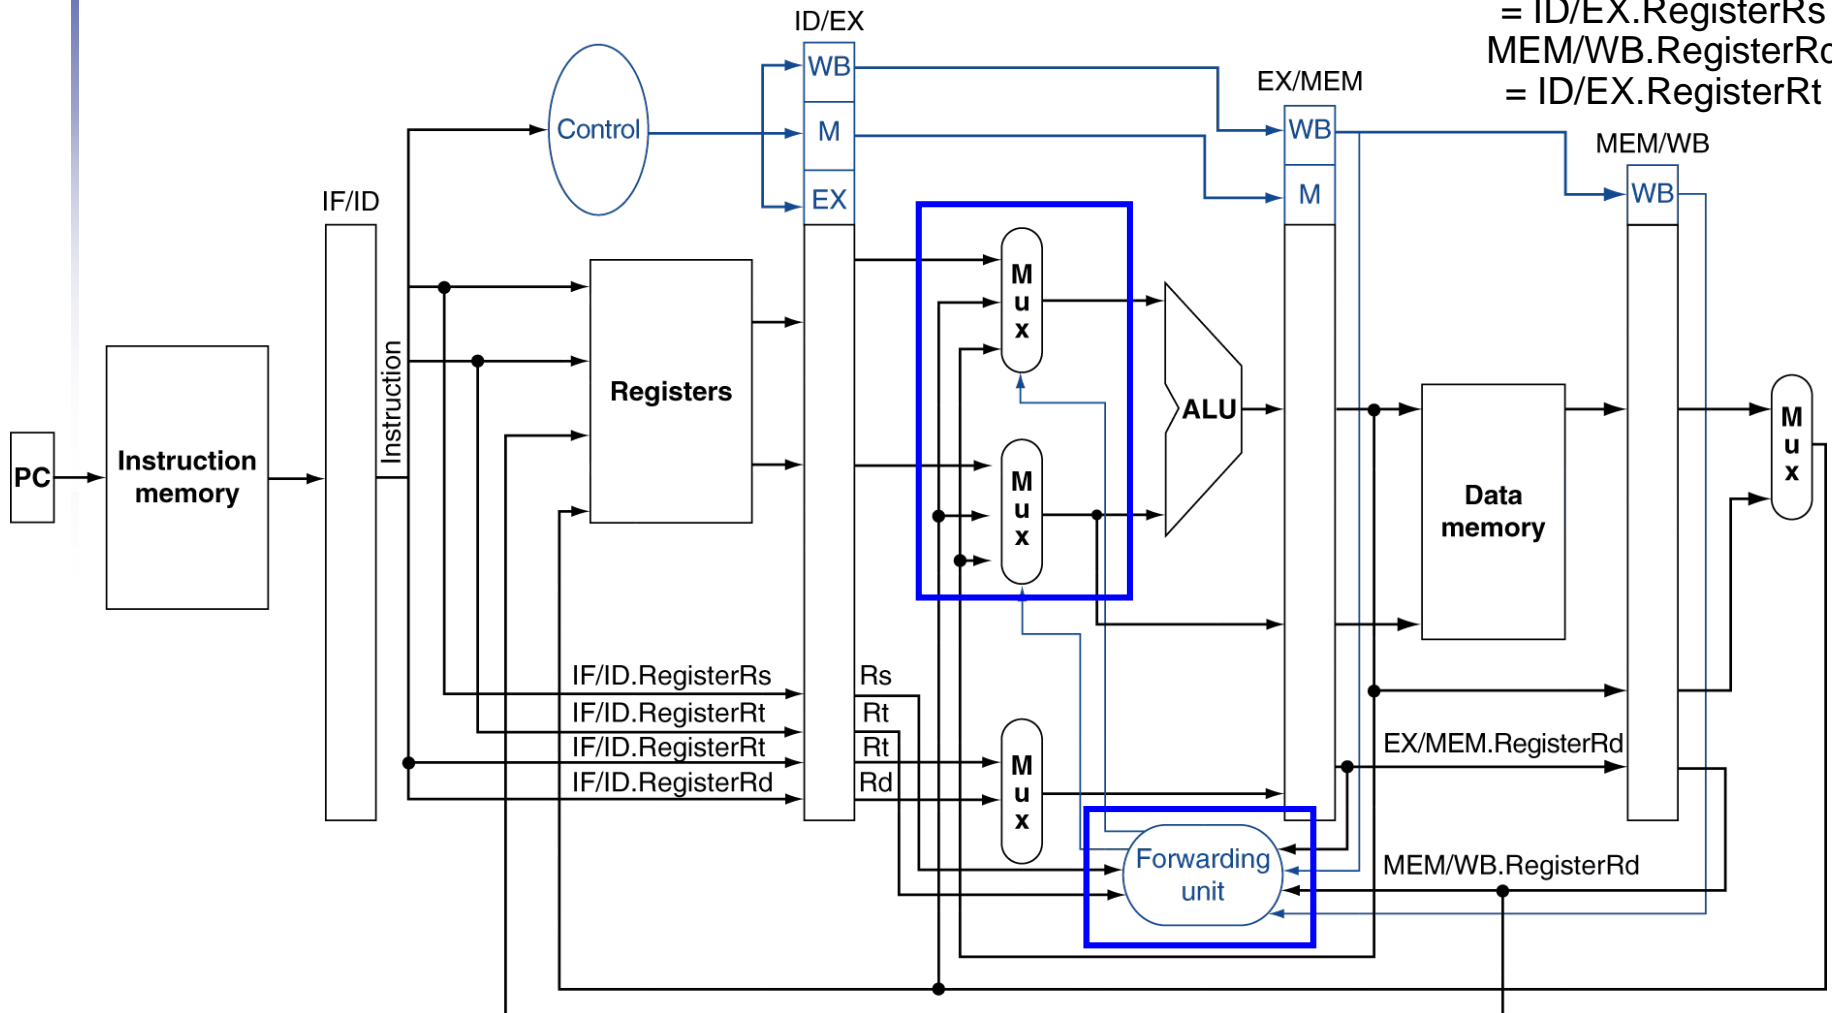
Fwd from  
MEM/WB  
pipeline reg

- *Example    Dependence Detection*

# Forwarding Logic : Fig. 4.56

EX/MEM.RegisterRd  
= ID/EX.RegisterRs  
EX/MEM.RegisterRd  
= ID/EX.RegisterRt

MEM/WB.RegisterRd  
= ID/EX.RegisterRs  
MEM/WB.RegisterRd  
= ID/EX.RegisterRt



# Forwarding Example

sub	\$2,	\$1,	\$3
and	\$4,	\$2,	\$5
or	\$4,	\$4,	\$2
add	\$9,	\$4,	\$2

# Forwarding Example

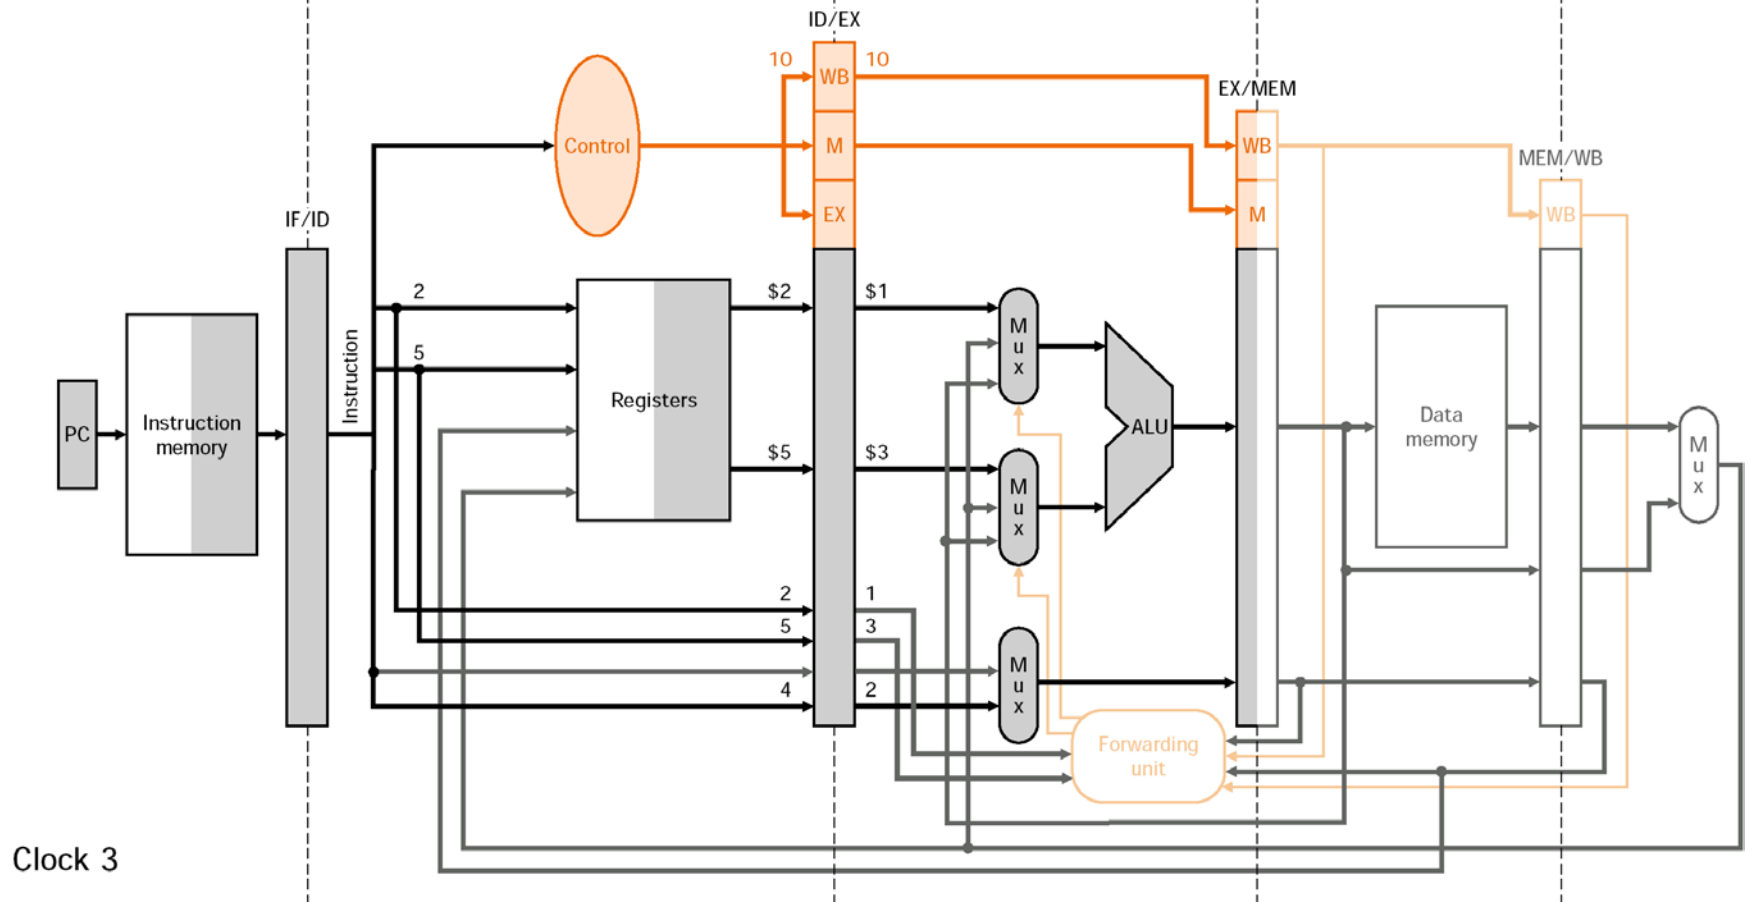
or \$4, \$4, \$2

and \$4, \$2, \$5

sub \$2, \$1, \$3

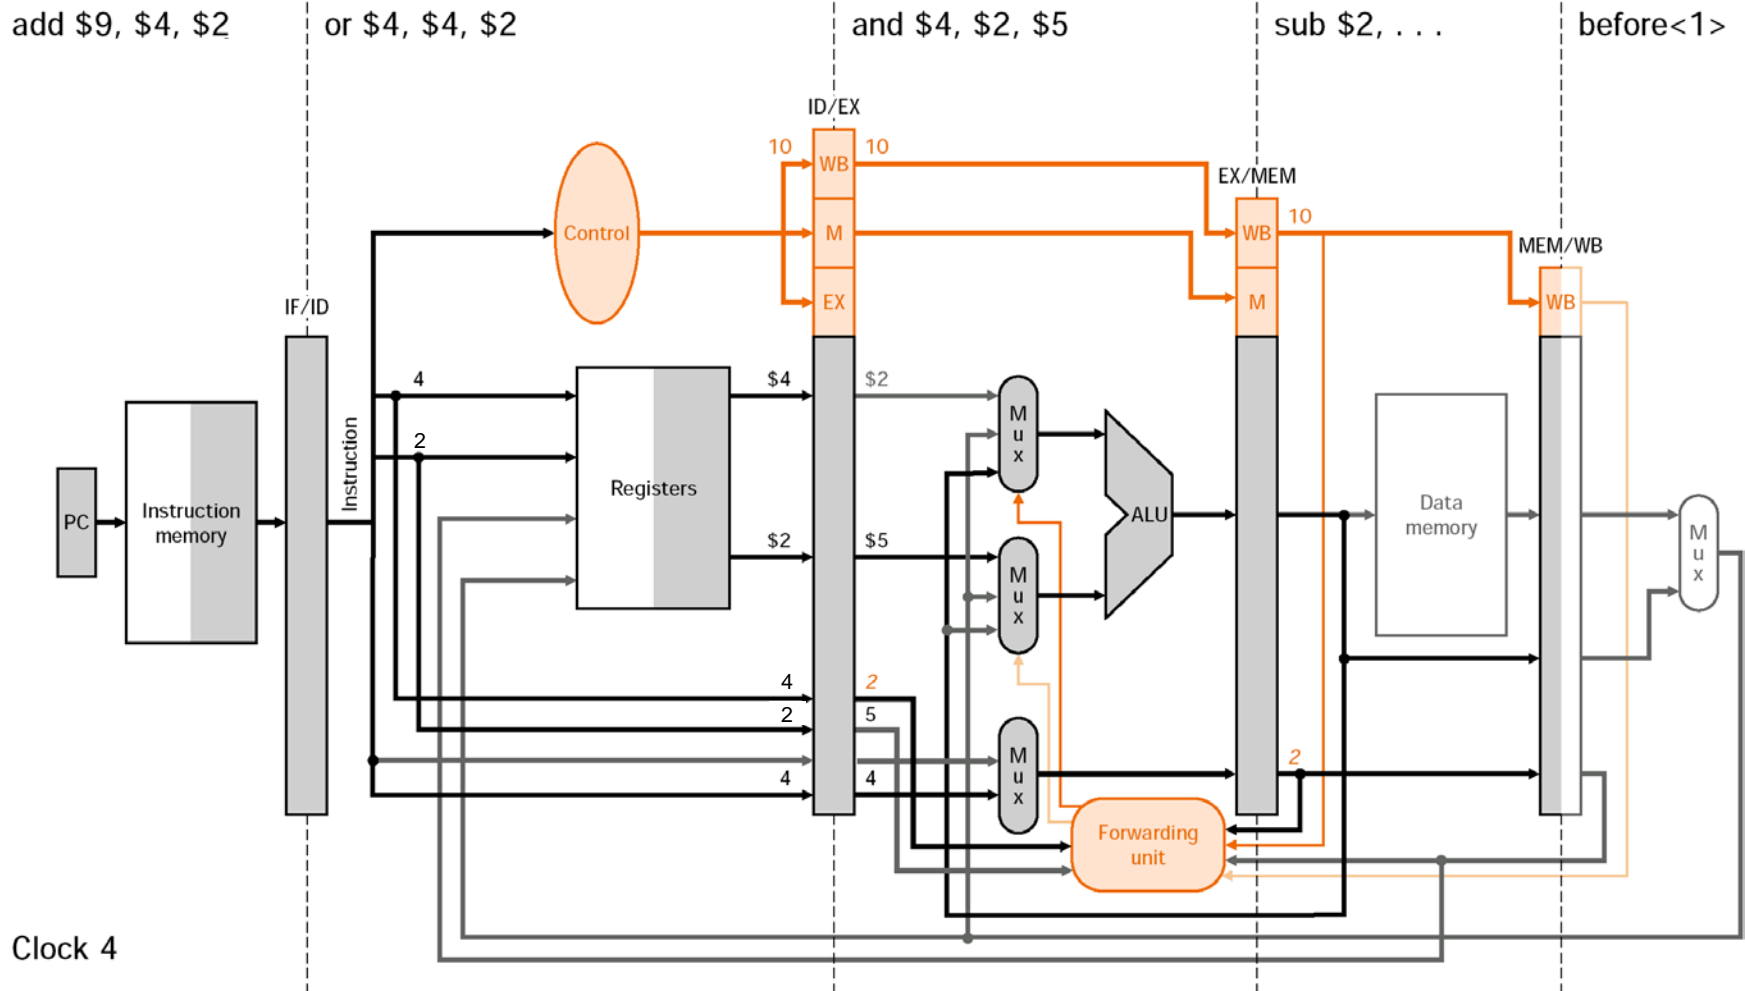
before<1>

before<2>



Clock 3

# Forwarding Example



# Forwarding Example


after<1>

add \$9, \$4, \$2

or \$4, \$4, \$2


and \$4, ...

sub \$2, ...



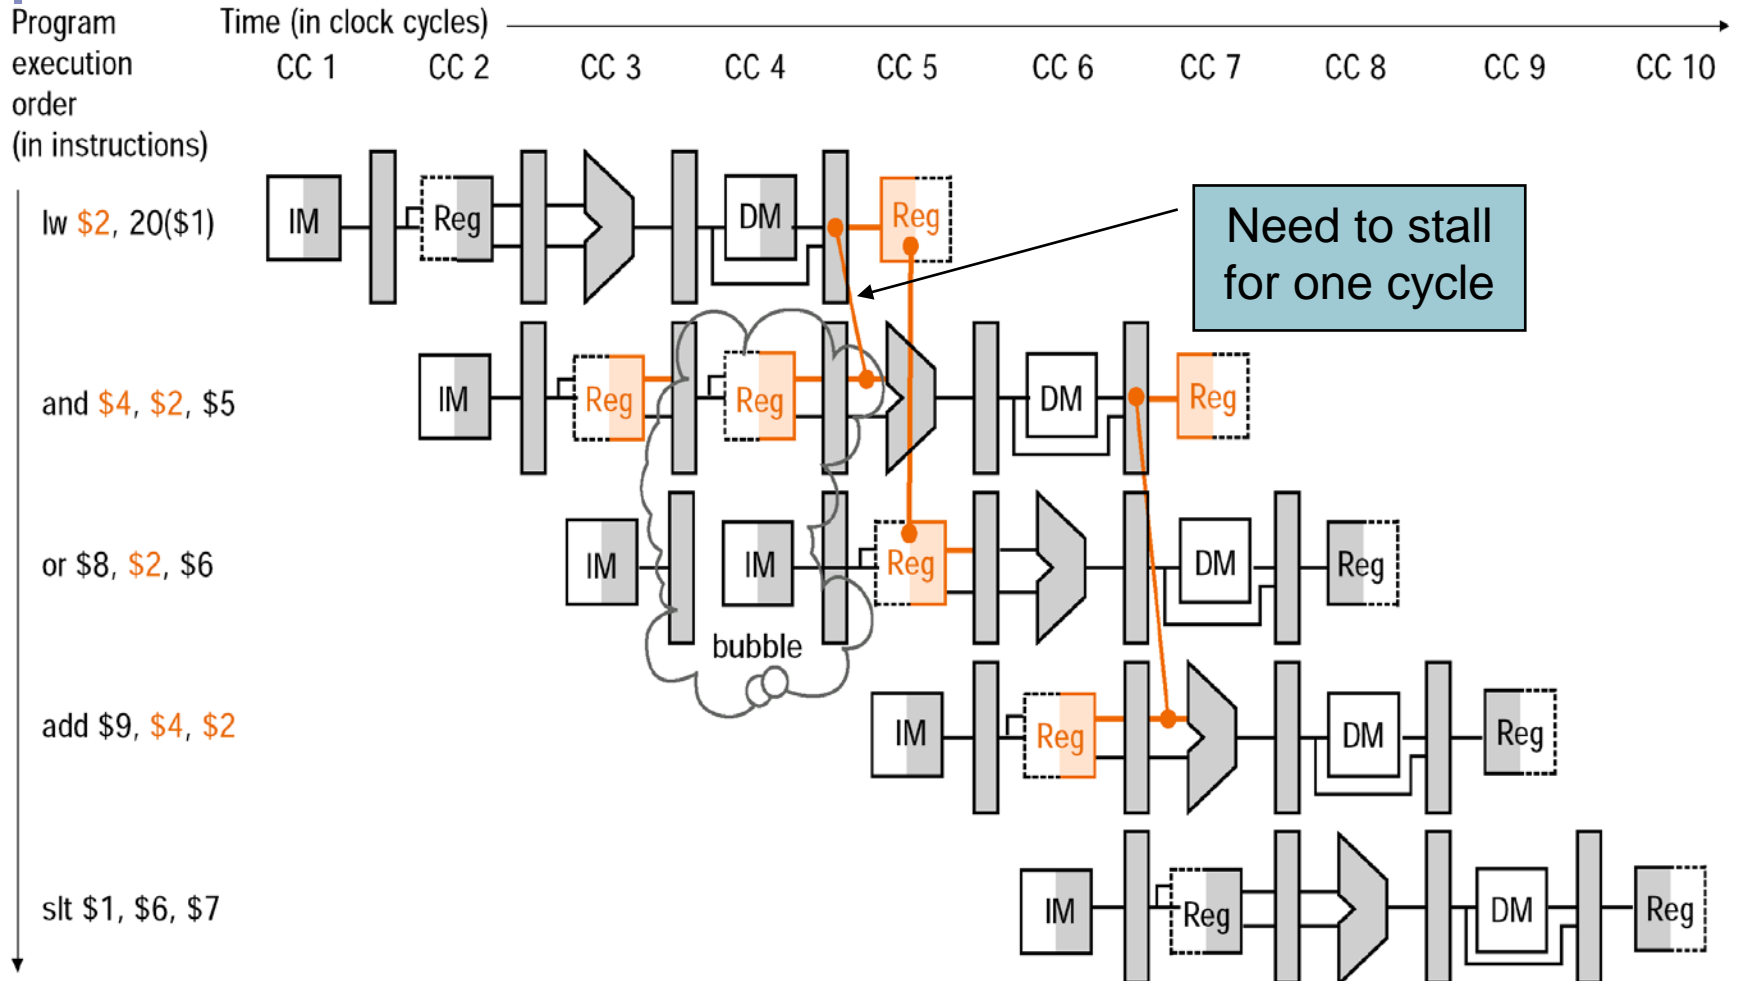


# Forwarding Example



# Data Hazard Requiring a Stall

Fig. 4.58

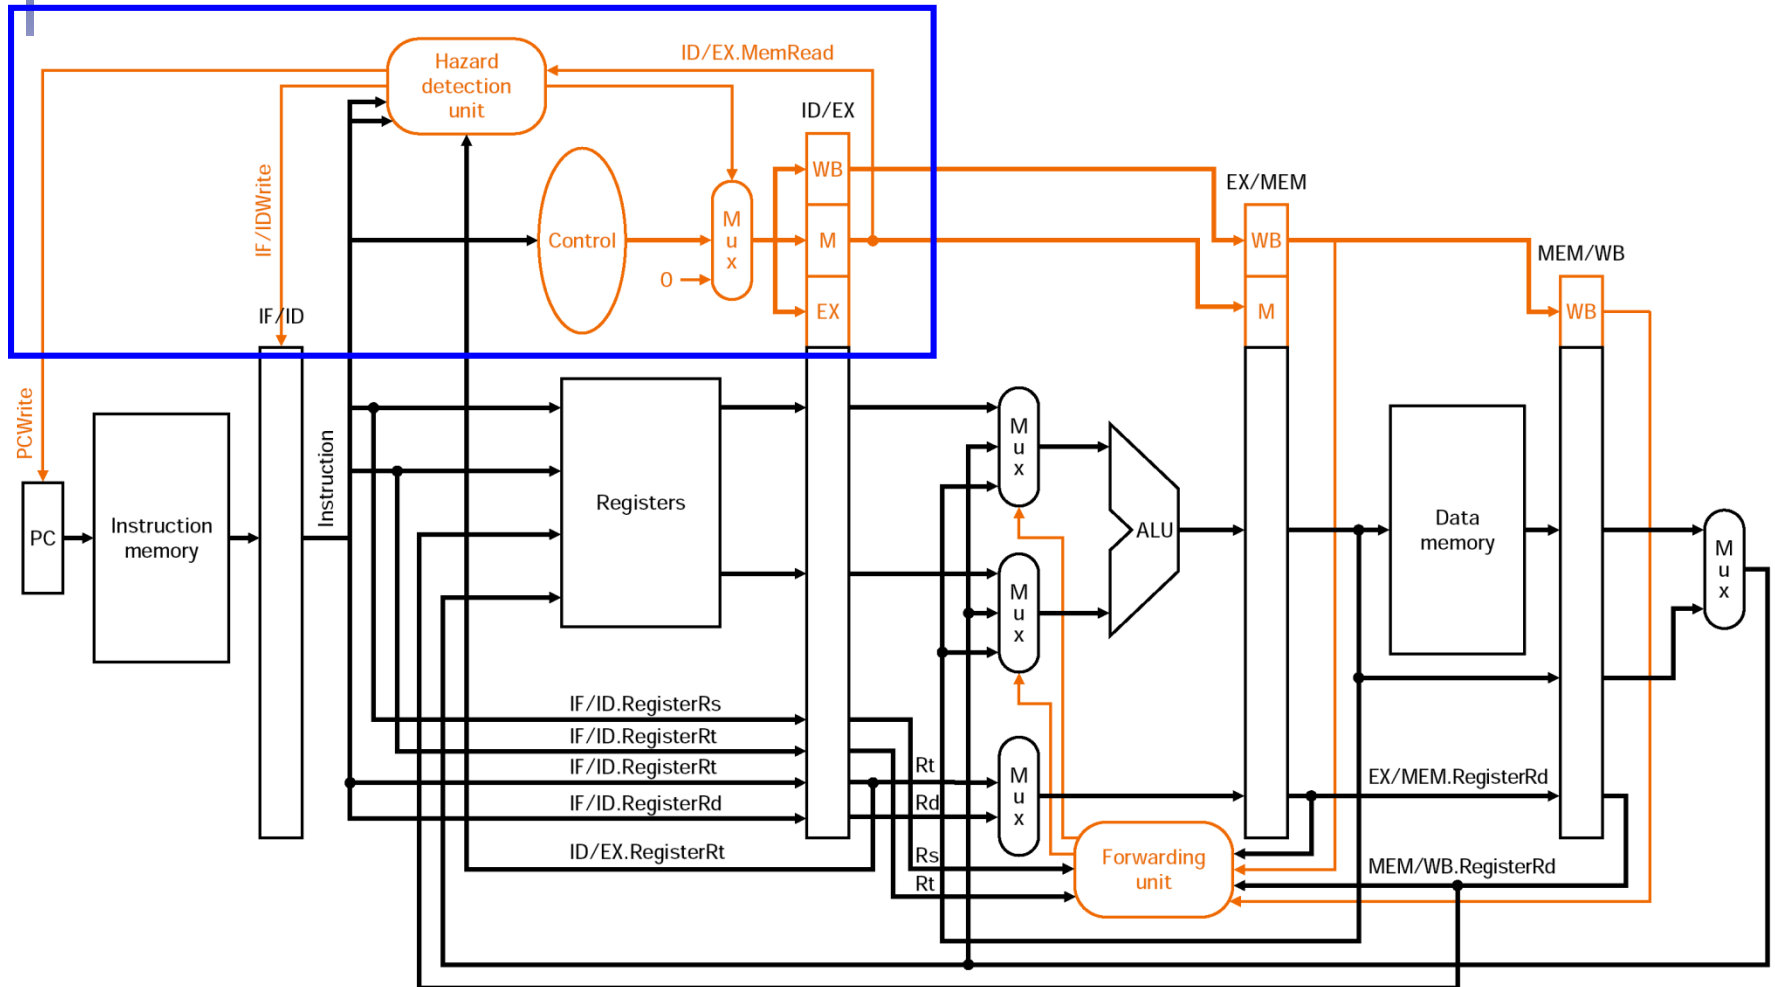


# Load-Use Hazard Detection

- Load-use hazard when
  - ID/EX.MemRead and  
((ID/EX.RegisterRt = IF/ID.RegisterRs) or  
(ID/EX.RegisterRt = IF/ID.RegisterRt))
- If detected, stall and insert bubble
- Force control values in ID/EX register to 0
  - EX, MEM and WB do nop (no-operation)
- Hazard Detection Unit

# Stall Logic

## Fig. 4.60



# Stall Example

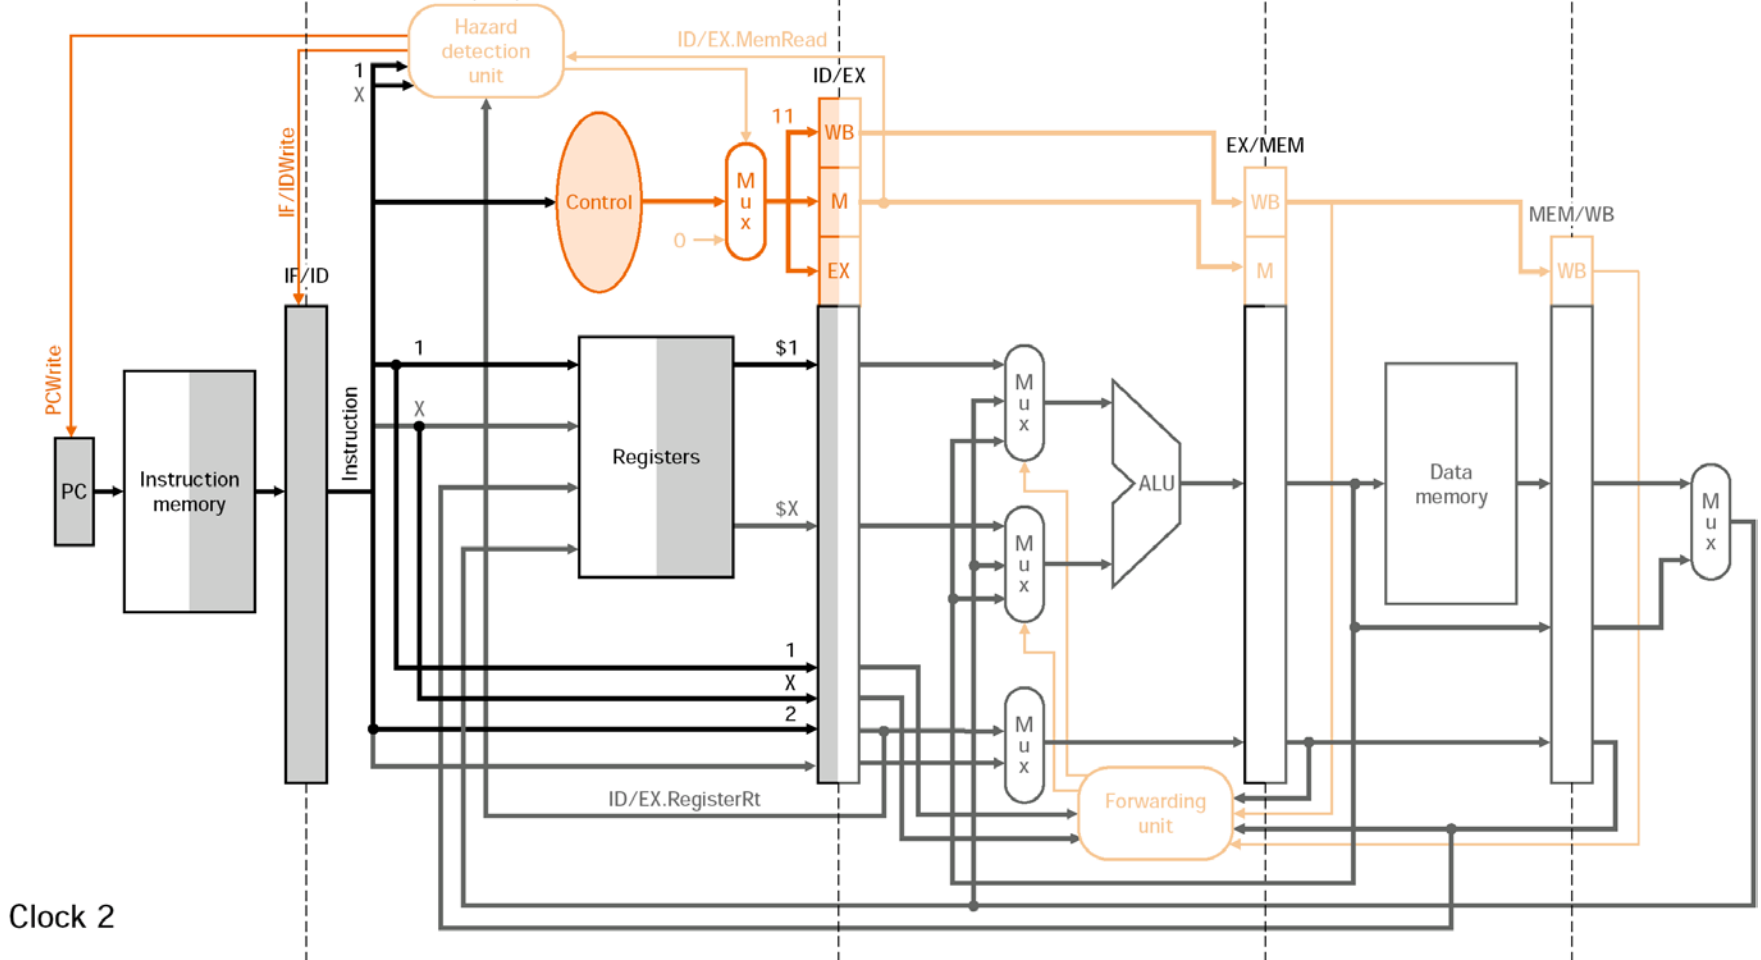
and \$4, \$2, \$5

lw \$2, 20(\$1)

before<1>

before<2>

before<3>



Clock 2

# Stall Example


or \$4, \$4, \$2

and \$4, \$2, \$5

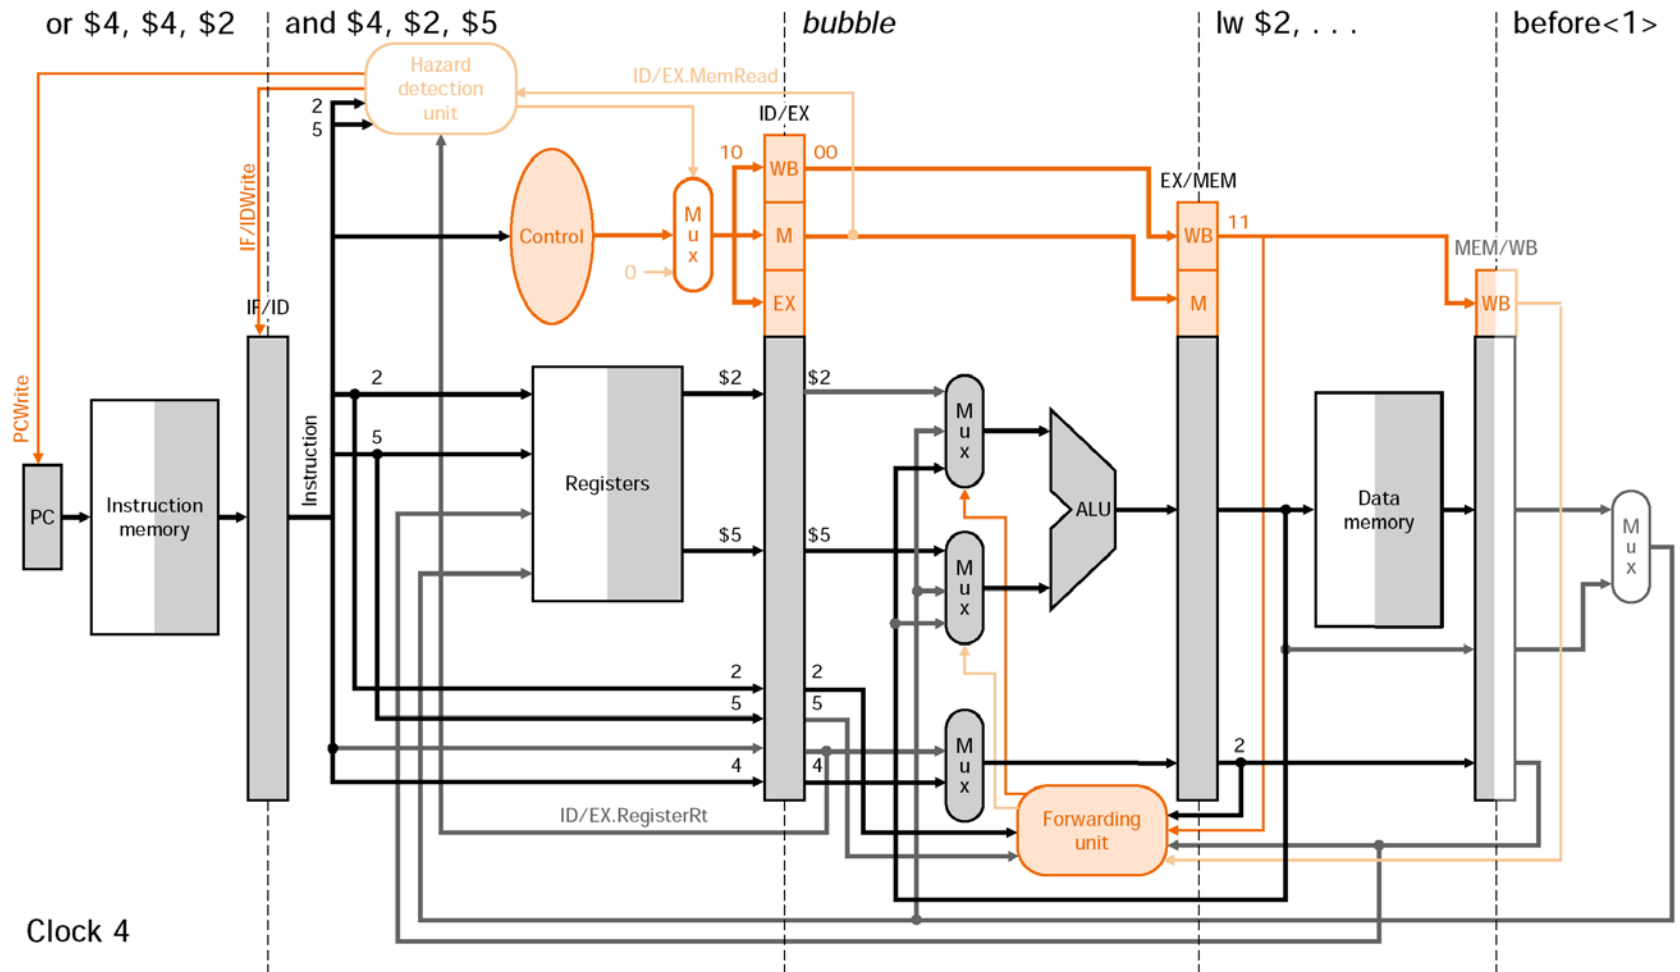
lw \$2, 20(\$1)

before<1>

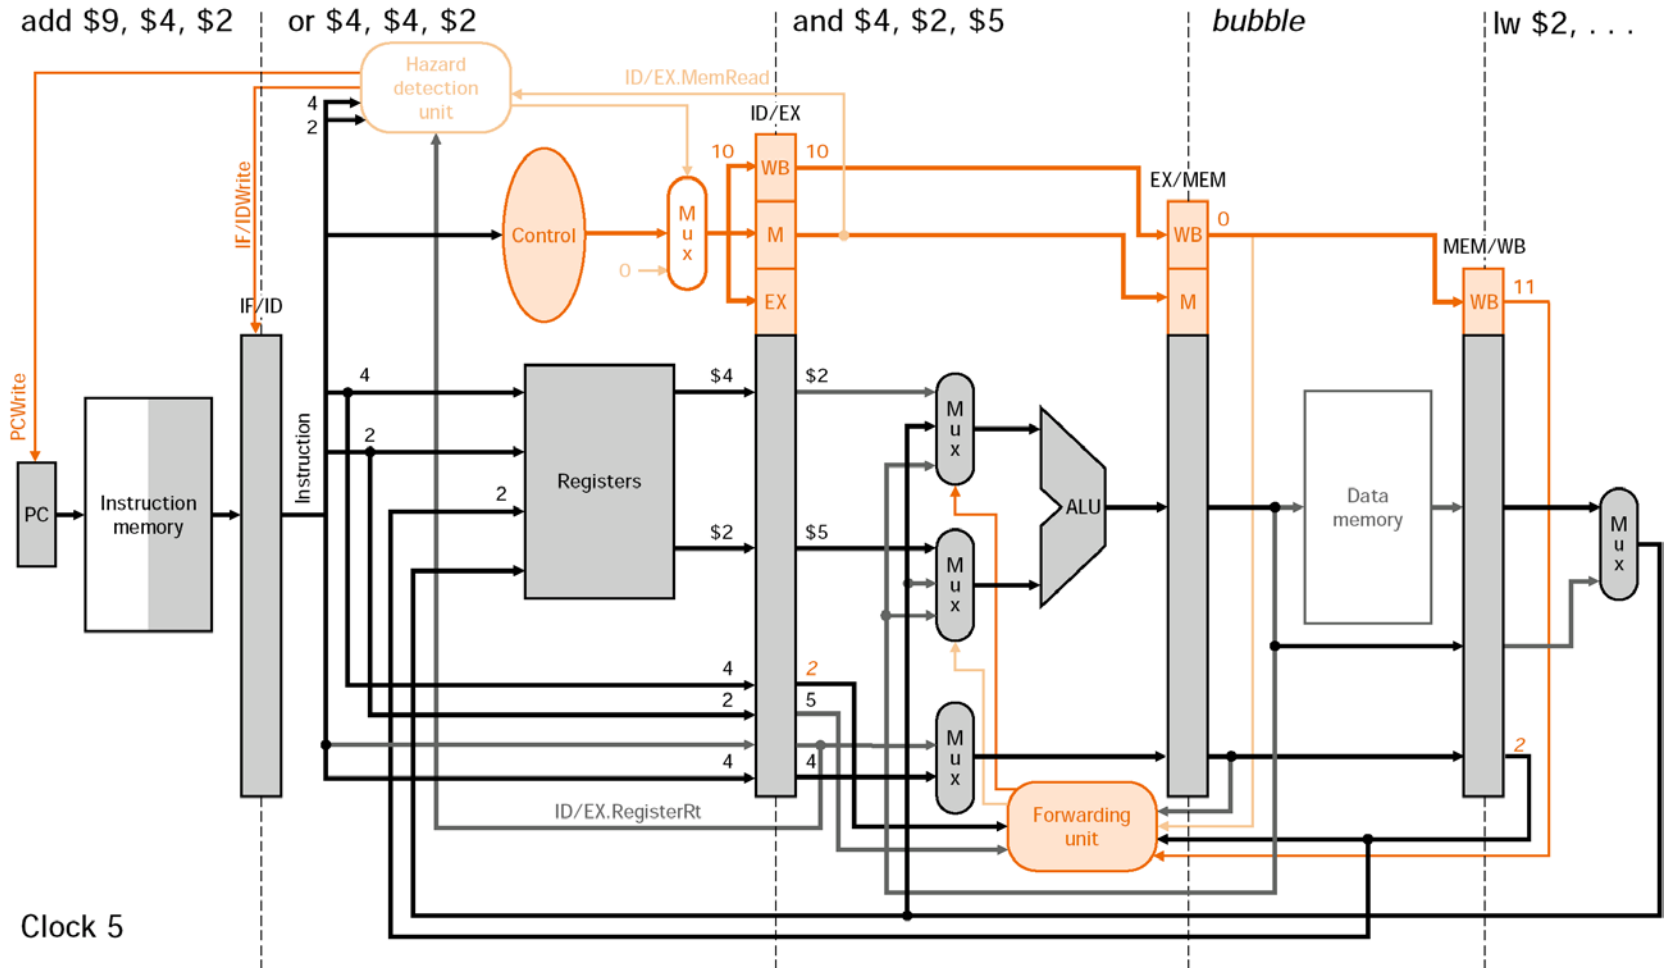
before<2>



# Stall Example



# Stall Example





# Stall Example

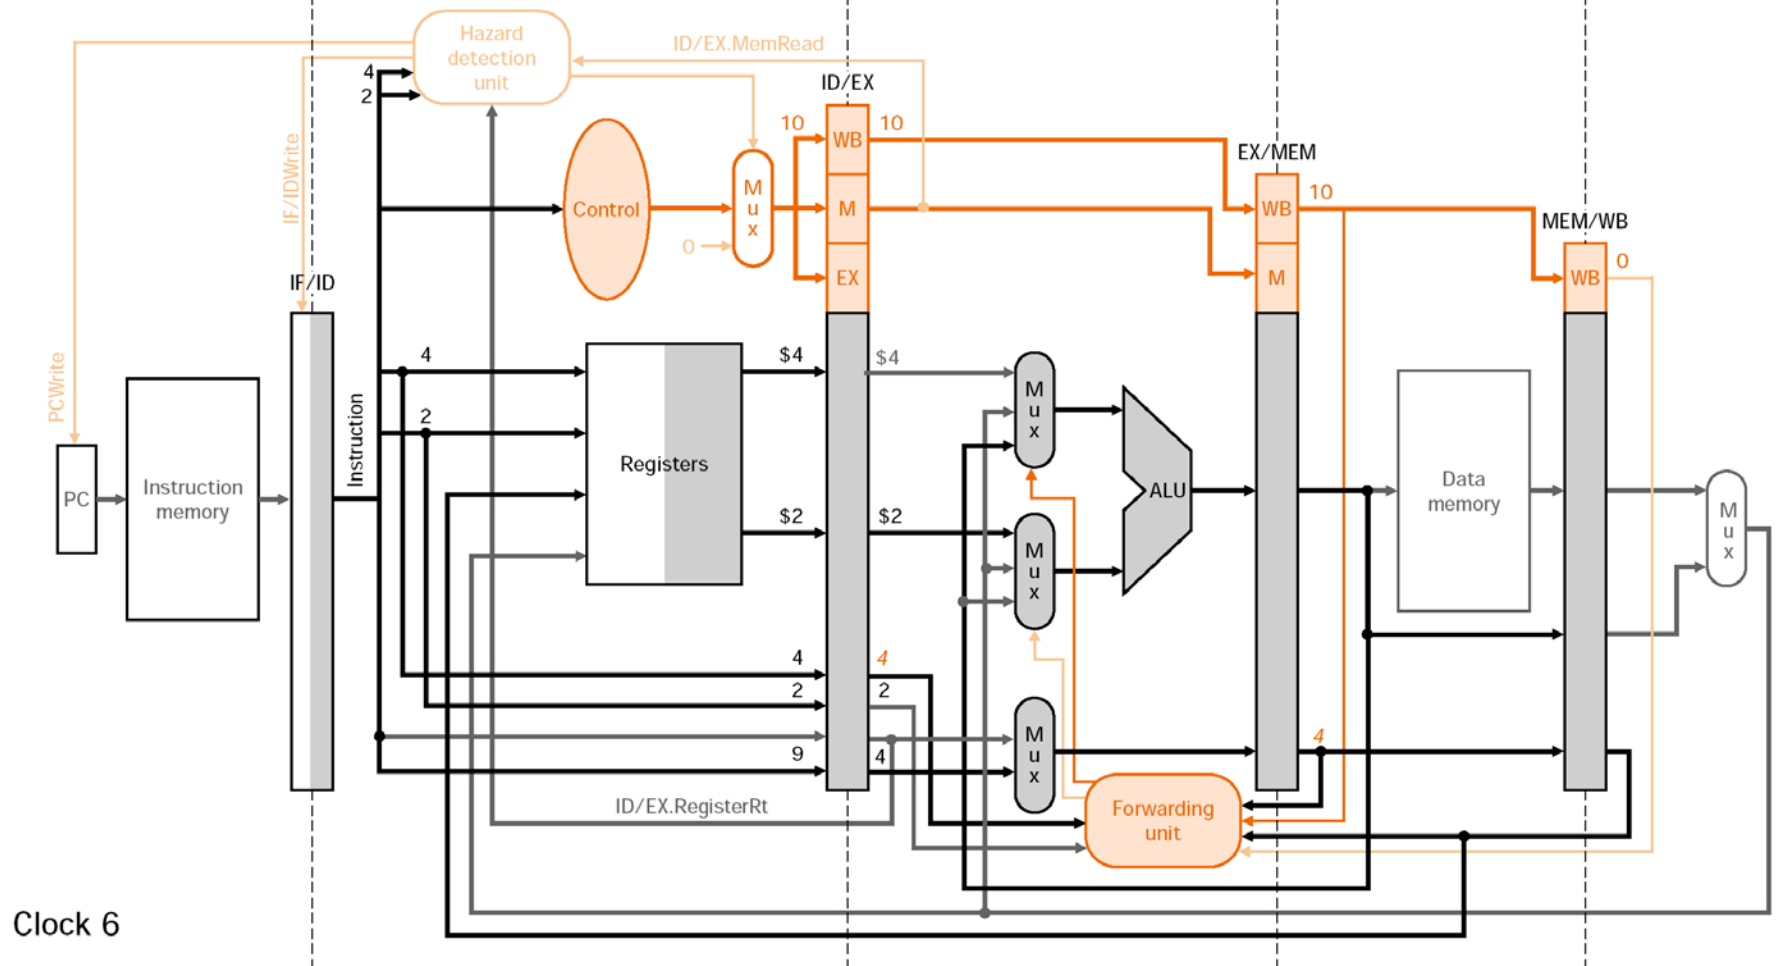
after<1>

add \$9, \$4, \$2

or \$4, \$4, \$2

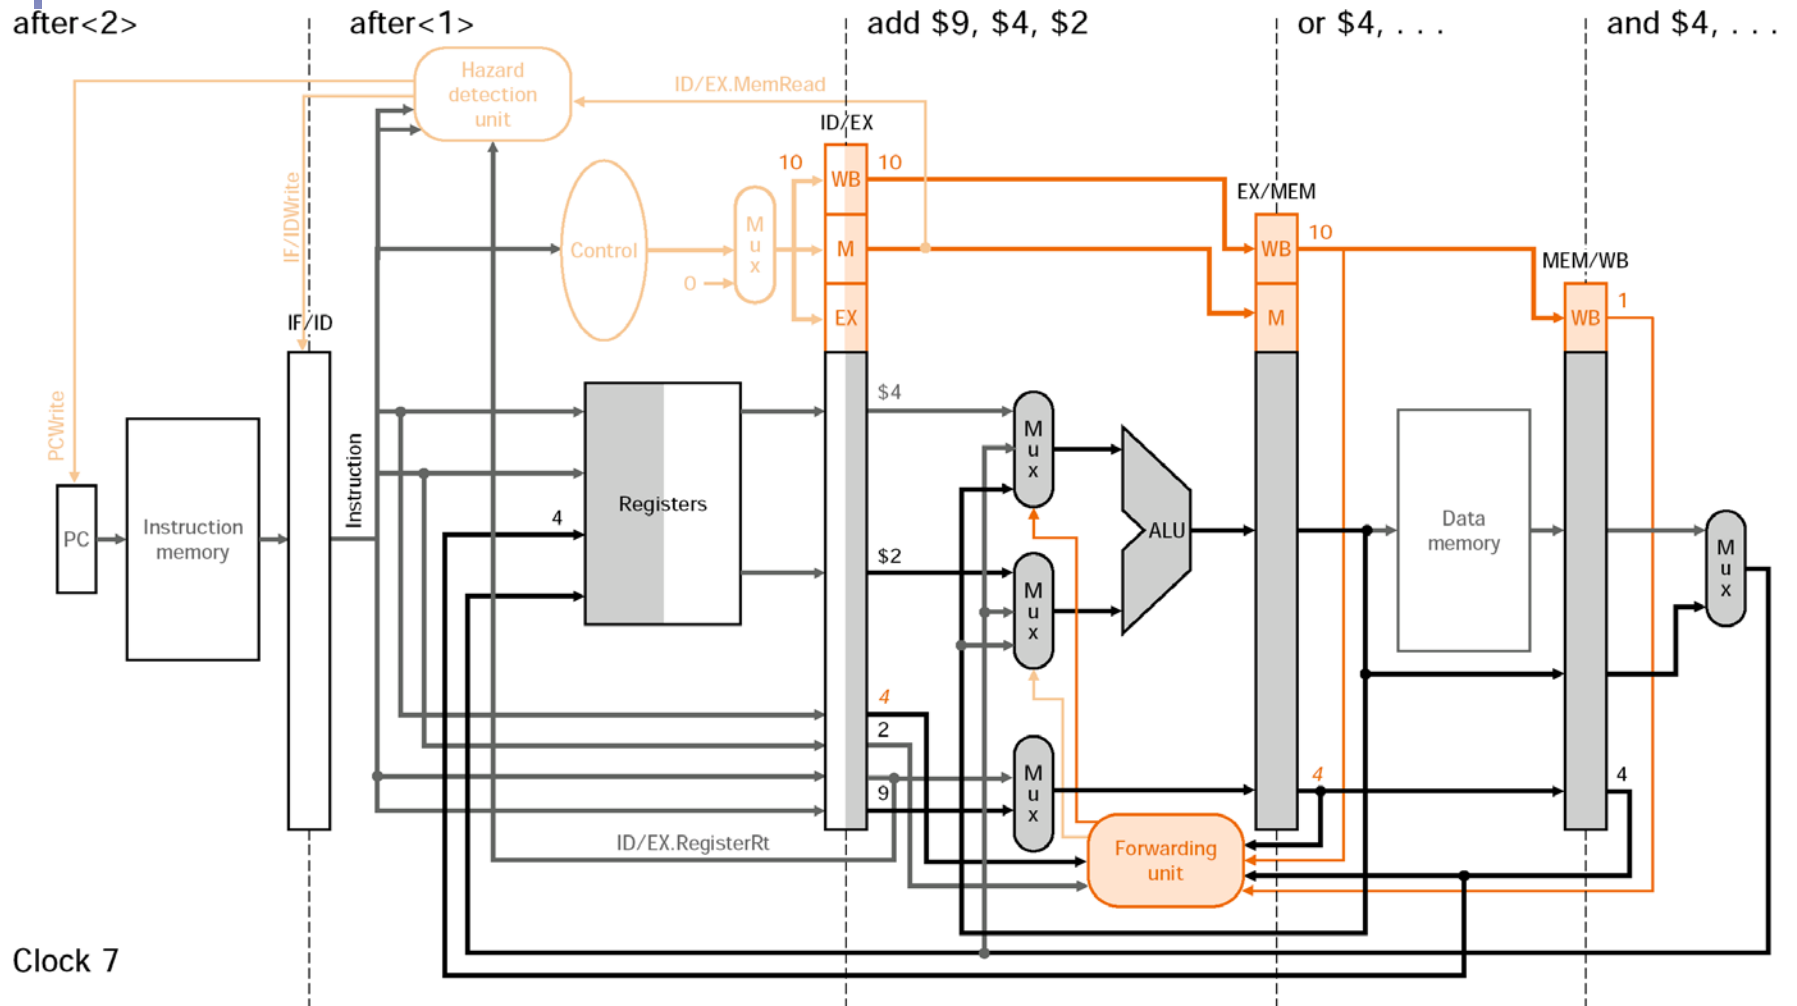
and \$4, ...

*bubble*



Clock 6

# Stall Example



Clock 7

# Stalls and Performance

- Stalls reduce performance
  - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
  - Requires knowledge of the pipeline structure

# Exercise #1: Without / With Forwarding

```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

	1	2	3	4	5	6	7	8	9	10	11
sub	IF	ID	EX	MEM	WB						
and		IF									
or											
add											
sw											

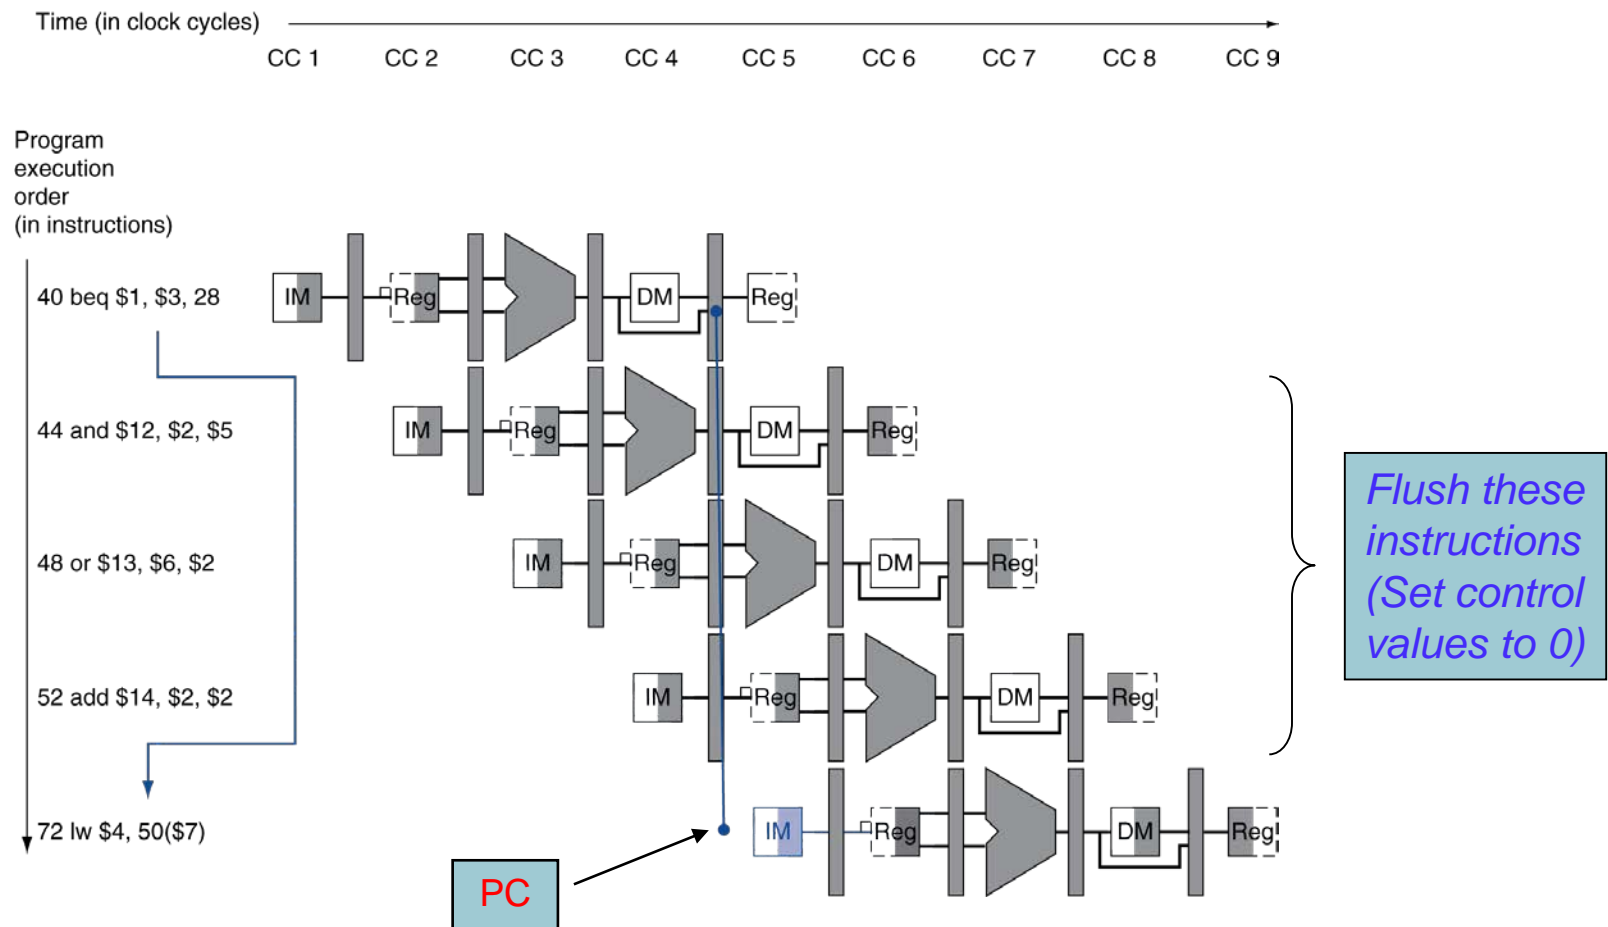
## Exercise #2: Without / With Forwarding

```
lw  $2, 20($3)
and $12, $2, $5
or  $13, $6, $2
add $14, $2, $2
sw  $15, 100($2)
```

	1	2	3	4	5	6	7	8	9	10	11
lw	IF	ID	EX	MEM	WB						
and		IF									
or											
add											
sw											

# Branch / Control Hazards

- If branch outcome determined in MEM Fig. 4.61



# Control Hazards: Reducing the Penalty

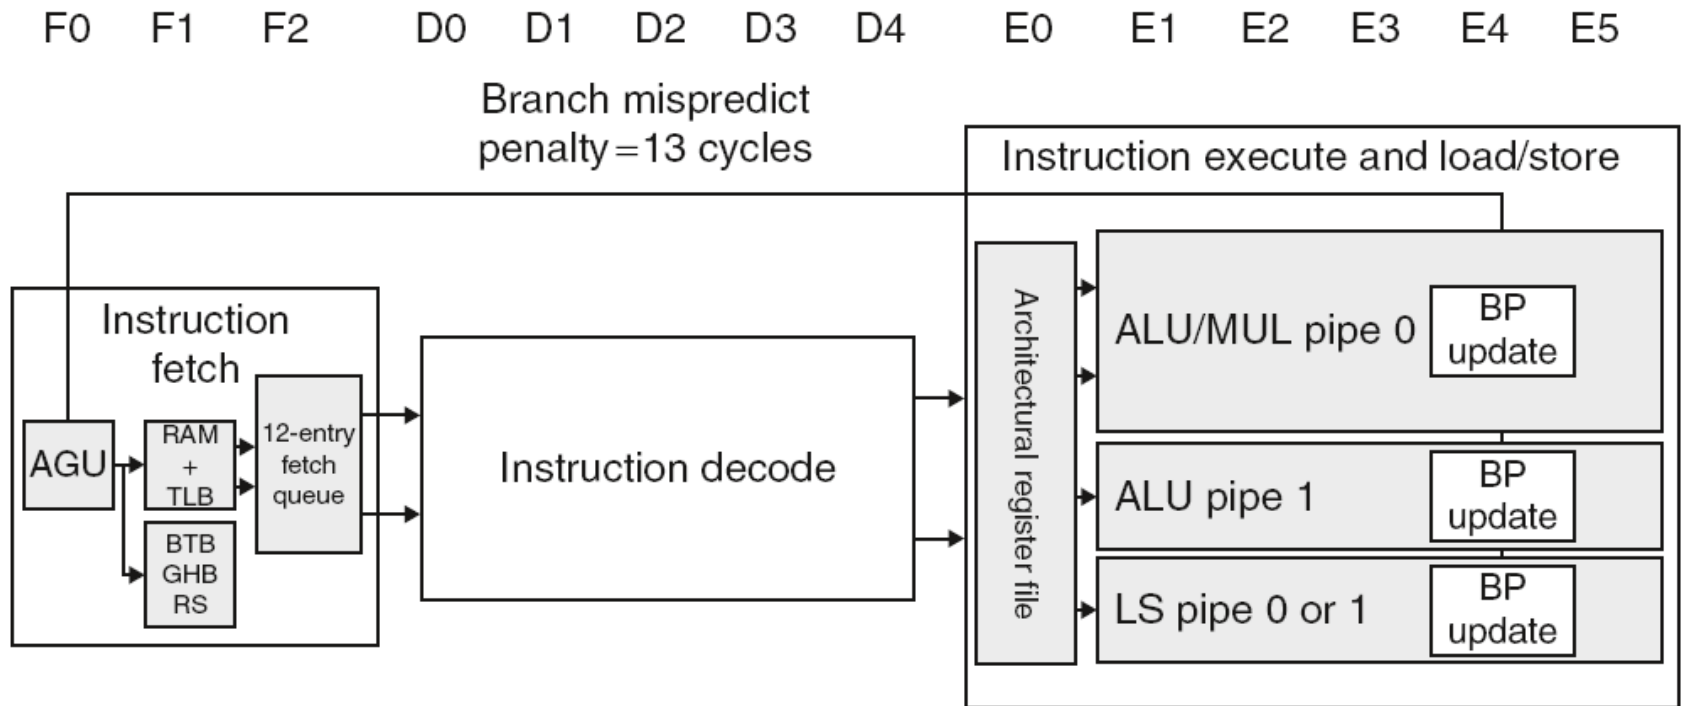
- Branching is very common in code:
  - A 3-cycle stall penalty is too heavy!
- Many techniques invented to reduce the control hazard penalty:
  - Move branch decision calculation to earlier pipeline stage
    - **Early Branch Resolution**
  - Guess the outcome before it is produced
    - **Branch Prediction**
  - Do something useful while waiting for the outcome
    - **Delayed Branching**

# Dynamic Branch Prediction

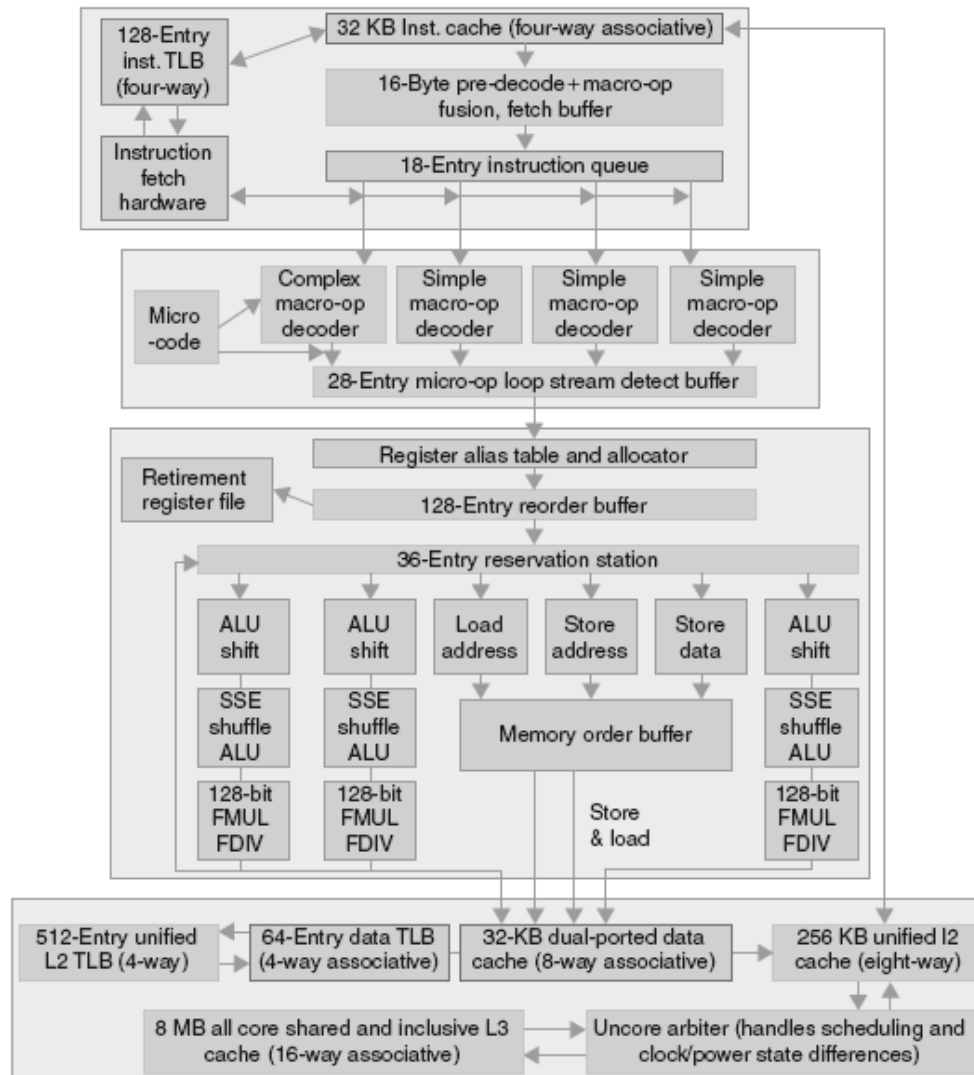
- In deeper and superscalar pipelines, branch penalty is more significant
  - *Reading pp. 318 – 319*
  - *Example pp. 319*
- Use **dynamic prediction**
  - Branch prediction buffer (aka branch history table)
  - Indexed by recent branch instruction addresses
  - Stores outcome (taken/not taken)
  - To execute a branch
    - Check table, expect the same outcome
    - Start fetching from fall-through or target
    - If wrong, flush pipeline and flip prediction



# ARM Cortex-A8 Pipeline



# Core i7 Pipeline



# Pipeline Summary

- ISA influences design of datapath and control
- Pipelining improves instruction throughput using parallelism
- Hazards : structural, data, control
- Implementation

## Chapter 4

Part D :

Exceptions

Parallelism via Instructions

# Exceptions and Interrupts

- “Unexpected” events requiring **change in flow of control**
  - Different ISAs use the terms differently
- Exception
  - Arises within the CPU
    - e.g., undefined opcode, overflow, syscall, ...
- Interrupt
  - From an external I/O controller
- Dealing with them without sacrificing performance is hard

# Handling Exceptions

- In MIPS, exceptions managed by a System Control **Coprocessor** (CP0)
- Save (E)PC of offending (or interrupted) instruction
- **Cause register vs. Vectored Interrupts**
  - Read cause, and transfer to relevant handler
    - Undefined opcode: C000 0000
    - Overflow: C000 0020
  - Determine action required
  - If restartable
    - Take corrective action
    - Use EPC to return to program
  - Otherwise
    - Terminate program
    - Report error using EPC, cause, ...

# Instruction-Level Parallelism (ILP)

- Pipelining : executing multiple instructions in parallel
- To increase ILP
  - Deeper pipeline
    - Less work per stage  $\Rightarrow$  shorter clock cycle
  - Multiple issue
    - Replicate pipeline stages  $\Rightarrow$  multiple pipelines
    - Start multiple instructions per clock cycle
    - But dependencies reduce this in practice

# Multiple Issue

- **Static** multiple issue
  - **Compiler** groups instructions to be issued together
  - Packages them into “issue slots”
  - Compiler detects and avoids hazards
- **Dynamic** multiple issue
  - **CPU** examines instruction stream and chooses instructions to issue each cycle
  - Compiler can help by reordering instructions
  - CPU resolves hazards using advanced techniques at runtime



# Static Multiple Issue

- Compiler groups instructions into “issue packets”
  - Group of instructions that can be issued on a single cycle
  - No dependencies within a packet
- Think of an issue packet as a very long instruction  
**Very Long Instruction Word (VLIW)**
  - Specifies multiple concurrent operations
  - Two-issue packets **Fig. 4.68, Example pp. 325**
    - One ALU/branch instruction
    - One load/store instruction

# Dynamic Multiple Issue

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, each cycle
  - Avoiding structural and data hazards
- Allow the CPU to execute instructions out of order to avoid stalls

- Example

```
lw      $t0, 20($s2)
addu    $t1, $t0, $t2
sub      $s4, $s4, $t3
slti    $t5, $s4, 20
```

- Can start sub while addu is waiting for lw

# Reading

- *pp. 342 ~ 344*

# Cortex A8 and Intel i7

Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level	2-level
1 <sup>st</sup> level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 <sup>nd</sup> level caches/core	128-1024 KiB	256 KiB
3 <sup>rd</sup> level caches (shared)	-	2- 8 MB

# Concluding Remarks

- ISA influences design of datapath and control
- Pipelining improves instruction throughput using parallelism
- Hazards : structural, data, control
- Multiple issue and dynamic scheduling (ILP)
  - Dependencies limit achievable parallelism
  - Complexity leads to the power wall